

```
In [1]: ▶ import pandas as pd
```

```
In [2]: ▶ df=pd.read_csv(r"C:\Users\Saiko\OneDrive\Desktop\Data Analytics\Car Price Pre
```

### Dataset Description:

The Car Price Dataset contains 10,000 records with 10 attributes detailing used cars and their resale prices. It includes brand, model, year (2000–2023), engine size (1.0L–5.0L), fuel type, transmission, mileage, doors, owner count, and price (2, 000–18,301). Newer cars, luxury brands, and lower mileage vehicles generally have higher prices. Automatic, diesel, and hybrid cars also tend to be more valuable. The dataset is ideal for price prediction models and market analysis, revealing trends such as depreciation patterns and the rising popularity of hybrid and electric vehicles due to environmental concerns.

### Columns in dataset:

- Brand (object): The car manufacturer (e.g., Kia, Chevrolet, Mercedes, Audi, etc.).
- Model (object): The specific model of the car.
- Year (int64): The manufacturing year of the car (range: 2000 to 2023).
- Engine\_Size (float64): The size of the engine in liters (range: 1.0L to 5.0L).
- Fuel\_Type (object): The type of fuel used (e.g., Diesel, Hybrid, Electric).
- Transmission (object): The type of transmission (e.g., Manual, Automatic, Semi-Automatic).
- Mileage (int64): The total distance the car has traveled, in kilometers (range: 25 to 299,947).
- Doors (int64): The number of doors (range: 2 to 5).
- Owner\_Count (int64): The number of previous owners (range: 1 to 5).
- Price (int64): The selling price of the car in USD (range: 2, 000 to 18,301)

In [3]: `df`

Out[3]:

	Brand	Model	Year	Engine_Size	Fuel_Type	Transmission	Mileage	Doors	Owner
0	Kia	Rio	2020	4.2	Diesel	Manual	289944	3	
1	Chevrolet	Malibu	2012	2.0	Hybrid	Automatic	5356	2	
2	Mercedes	GLA	2020	4.2	Diesel	Automatic	231440	4	
3	Audi	Q5	2023	2.0	Electric	Manual	160971	2	
4	Volkswagen	Golf	2003	2.6	Hybrid	Semi-Automatic	286618	3	
...	...	...	...	...	...	...	...	...	...
9995	Kia	Optima	2004	3.7	Diesel	Semi-Automatic	5794	2	
9996	Chevrolet	Impala	2002	1.4	Electric	Automatic	168000	2	
9997	BMW	3 Series	2010	3.0	Petrol	Automatic	86664	5	
9998	Ford	Explorer	2002	1.4	Hybrid	Automatic	225772	4	
9999	Volkswagen	Tiguan	2001	2.1	Diesel	Manual	157882	3	

10000 rows × 10 columns

In [4]: `df.isnull().sum()`

```
Out[4]: Brand      0
Model      0
Year       0
Engine_Size 0
Fuel_Type  0
Transmission 0
Mileage    0
Doors      0
Owner_Count 0
Price      0
dtype: int64
```

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 10 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   Brand           10000 non-null  object  
 1   Model           10000 non-null  object  
 2   Year            10000 non-null  int64   
 3   Engine_Size     10000 non-null  float64  
 4   Fuel_Type       10000 non-null  object  
 5   Transmission    10000 non-null  object  
 6   Mileage         10000 non-null  int64   
 7   Doors           10000 non-null  int64   
 8   Owner_Count     10000 non-null  int64   
 9   Price           10000 non-null  int64   
dtypes: float64(1), int64(5), object(4)
memory usage: 781.4+ KB
```

In [6]: `from sklearn.linear_model import LinearRegression`  
`from sklearn import linear_model`

In [7]: `brand=pd.get_dummies(df['Brand'],prefix="Brand")`

In [8]: `model=pd.get_dummies(df['Model'],prefix="Model")`

In [9]: `fuel_type=pd.get_dummies(df['Fuel_Type'],prefix="Fuel_Type")`

In [10]: `transmission=pd.get_dummies(df['Transmission'],prefix="Transmission")`

In [11]: `print(brand)`

	Brand_Audi	Brand_BMW	Brand_Chevrolet	Brand_Ford	Brand_Honda	\
0	False	False	False	False	False	
1	False	False	True	False	False	
2	False	False	False	False	False	
3	True	False	False	False	False	
4	False	False	False	False	False	
...	...	...	...	...	...	
9995	False	False	False	False	False	
9996	False	False	True	False	False	
9997	False	True	False	False	False	
9998	False	False	False	True	False	
9999	False	False	False	False	False	
	Brand_Hyundai	Brand_Kia	Brand_Mercedes	Brand_Toyota	Brand_Volkswa	
gen						
0	False	True	False	False		Fa
lse						
1	False	False	False	False		Fa
lse						
2	False	False	True	False		Fa
lse						
3	False	False	False	False		Fa
lse						
4	False	False	False	False		T
rue						
...	...	...	...	...		
...						
9995	False	True	False	False		Fa
lse						
9996	False	False	False	False		Fa
lse						
9997	False	False	False	False		Fa
lse						
9998	False	False	False	False		Fa
lse						
9999	False	False	False	False		T
rue						

[10000 rows x 10 columns]

In [12]: ▶ `print(model)`

	Model_3	Series	Model_5	Series	Model_A3	Model_A4	Model_Accord	\
0		False		False	False	False	False	
1		False		False	False	False	False	
2		False		False	False	False	False	
3		False		False	False	False	False	
4		False		False	False	False	False	
...		...		...	...	...	...	
9995		False		False	False	False	False	
9996		False		False	False	False	False	
9997		True		False	False	False	False	
9998		False		False	False	False	False	
9999		False		False	False	False	False	

	Model_C-Class	Model_CR-V	Model_Camry	Model_Civic	Model_Corolla	\
...						
0	False	False	False	False	False	
...						
1	False	False	False	False	False	
...						
2	False	False	False	False	False	
...						
3	False	False	False	False	False	
...						
4	False	False	False	False	False	
...						
...	...	...	...	...	...	
...						
9995	False	False	False	False	False	
...						
9996	False	False	False	False	False	
...						
9997	False	False	False	False	False	
...						
9998	False	False	False	False	False	
...						
9999	False	False	False	False	False	
...						

	Model_Optima	Model_Passat	Model_Q5	Model_RAV4	Model_Rio	\
0	False	False	False	False	True	
1	False	False	False	False	False	
2	False	False	False	False	False	
3	False	False	True	False	False	
4	False	False	False	False	False	
...	...	...	...	...	...	
9995	True	False	False	False	False	
9996	False	False	False	False	False	
9997	False	False	False	False	False	
9998	False	False	False	False	False	
9999	False	False	False	False	False	

	Model_Sonata	Model_Sportage	Model_Tiguan	Model_Tucson	Model_X5
0	False	False	False	False	False
1	False	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False

```
...
9995      False      False      False      False      False
9996      False      False      False      False      False
9997      False      False      False      False      False
9998      False      False      False      False      False
9999      False      False      True       False      False
```

[10000 rows x 30 columns]

In [13]:

▶ print(transmission)

```
Transmission_Automatic  Transmission_Manual  Transmission_Semi-Automa
tic
0      False      True      Fa
lse
1      True      False      Fa
lse
2      True      False      Fa
lse
3      False      True      Fa
lse
4      False      False     T
rue
...      ...      ...
...
9995     False     False     T
rue
9996     True      False     Fa
lse
9997     True      False     Fa
lse
9998     True      False     Fa
lse
9999     False     True      Fa
lse
```

[10000 rows x 3 columns]

In [14]: `print(fuel_type)`

	Fuel_Type_Diesel	Fuel_Type_Electric	Fuel_Type_Hybrid	Fuel_Type_Pet
rol				
0	True	False	False	Fa
lse				
1	False	False	True	Fa
lse				
2	True	False	False	Fa
lse				
3	False	True	False	Fa
lse				
4	False	False	True	Fa
lse				
...	...	...	...	
...				
9995	True	False	False	Fa
lse				
9996	False	True	False	Fa
lse				
9997	False	False	False	T
rue				
9998	False	False	True	Fa
lse				
9999	True	False	False	Fa
lse				

[10000 rows x 4 columns]

In [15]: `df.drop(["Brand", "Model", "Fuel_Type", "Transmission"], axis=1, inplace=True)`

In [16]: `df=pd.concat([df,brand,model,fuel_type,transmission],axis=1)`

In [17]: `reg=linear_model.LinearRegression()  
reg.fit(df.drop('Price',axis='columns'),df.Price)`

Out[17]: LinearRegression()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**



In [18]: `reg.coef_`

```
Out[18]: array([ 2.98601356e+02,  9.92739311e+02, -1.98902437e-02, -5.50438898e-01,
  3.58753313e-02, -3.00390404e+00, -9.13203152e-01, -1.55918923e+00,
 -1.38937007e-01,  3.95555544e+00, -1.86013447e+00,  4.31507789e+00,
  2.03848302e+00, -1.10822340e+00, -1.72552506e+00,  2.86758378e-01,
 -6.77909298e+00,  1.13687704e+00, -3.18040004e+00, -9.45409685e-01,
 -6.93412569e-01, -1.24997283e+00, -4.99892135e+00,  6.15093795e+00,
 -2.03858422e+00, -3.47440863e+00, -2.29589683e+00, -3.14473413e+00,
  3.99551403e-01, -1.85058540e+00,  1.31209699e+00,  6.20630422e+00,
 -8.45653658e-01,  1.79503760e+00, -2.09492702e-01,  1.22018606e+00,
 -1.77553804e-02, -9.60381043e-01,  5.92928217e+00,  1.26855132e+00,
 -1.08576609e+00,  1.82634052e+00, -8.62116017e-01,  1.52152846e+00,
  5.57913145e+00, -7.44851446e+02,  1.24307000e+03,  2.45486884e+02,
 -7.43705441e+02,  9.94064951e+02, -4.96166920e+02, -4.97898031e+02])
```

In [19]: `reg.intercept_`

```
Out[19]: -591822.953814718
```

In [20]: `df`

Out[20]:

	Year	Engine_Size	Mileage	Doors	Owner_Count	Price	Brand_Audi	Brand_BMW	Brai
0	2020	4.2	289944	3	5	8501	False	False	
1	2012	2.0	5356	2	3	12092	False	False	
2	2020	4.2	231440	4	2	11171	False	False	
3	2023	2.0	160971	2	1	11780	True	False	
4	2003	2.6	286618	3	3	2867	False	False	
...	...	...	...	...	...	...	...	...	
9995	2004	3.7	5794	2	4	8884	False	False	
9996	2002	1.4	168000	2	1	6240	False	False	
9997	2010	3.0	86664	5	1	9866	False	True	
9998	2002	1.4	225772	4	1	4084	False	False	
9999	2001	2.1	157882	3	3	3342	False	False	

10000 rows × 53 columns



In [21]: `import matplotlib.pyplot as plt`  
`import seaborn as sns`


In [22]: `plt.figure(figsize=(15, 10))`

```
Out[22]: <Figure size 1500x1000 with 0 Axes>
```

```
<Figure size 1500x1000 with 0 Axes>
```

```
In [23]: sns.histplot(df['Price'], kde=True, color='skyblue')  
plt.title('Distribution of Car Prices')
```

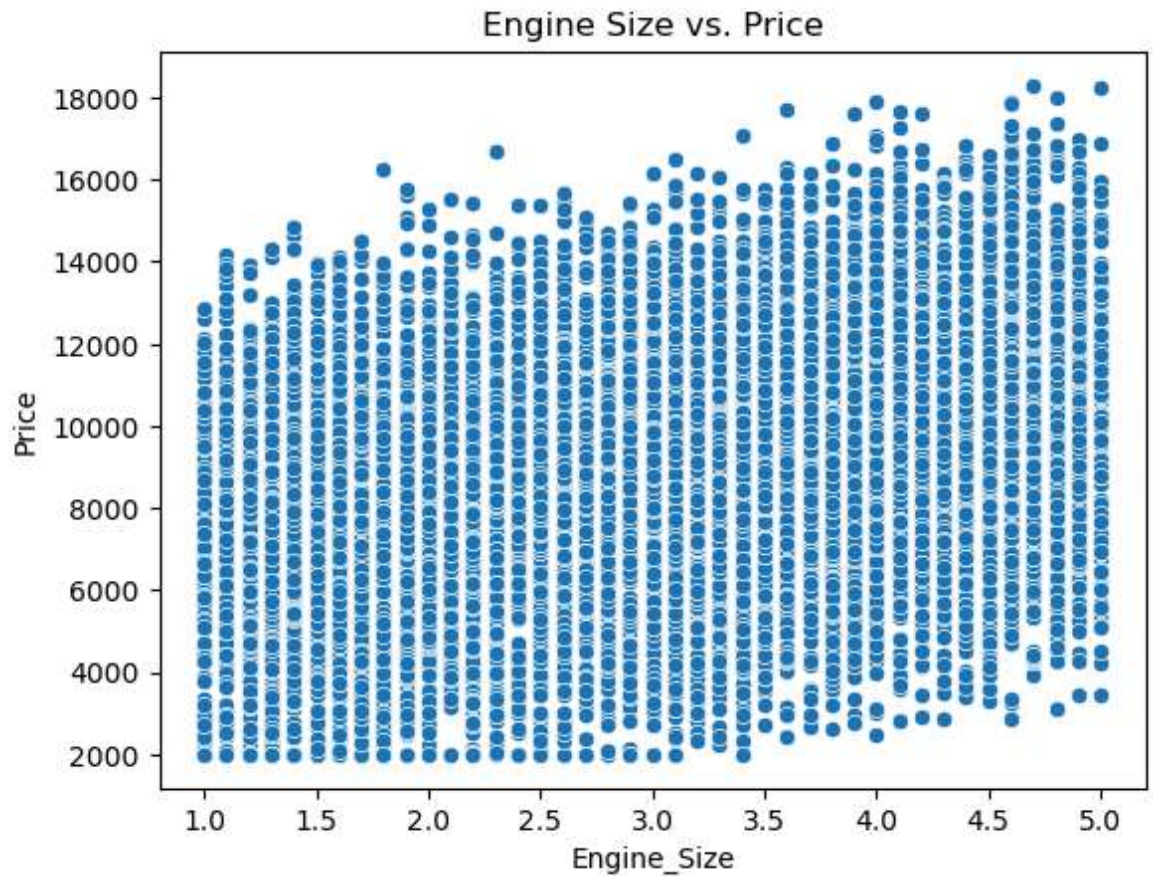
Out[23]: Text(0.5, 1.0, 'Distribution of Car Prices')



```
In [24]: sns.scatterplot(x='Engine_Size', y='Price', data=df, palette='viridis')  
plt.title('Engine Size vs. Price')
```

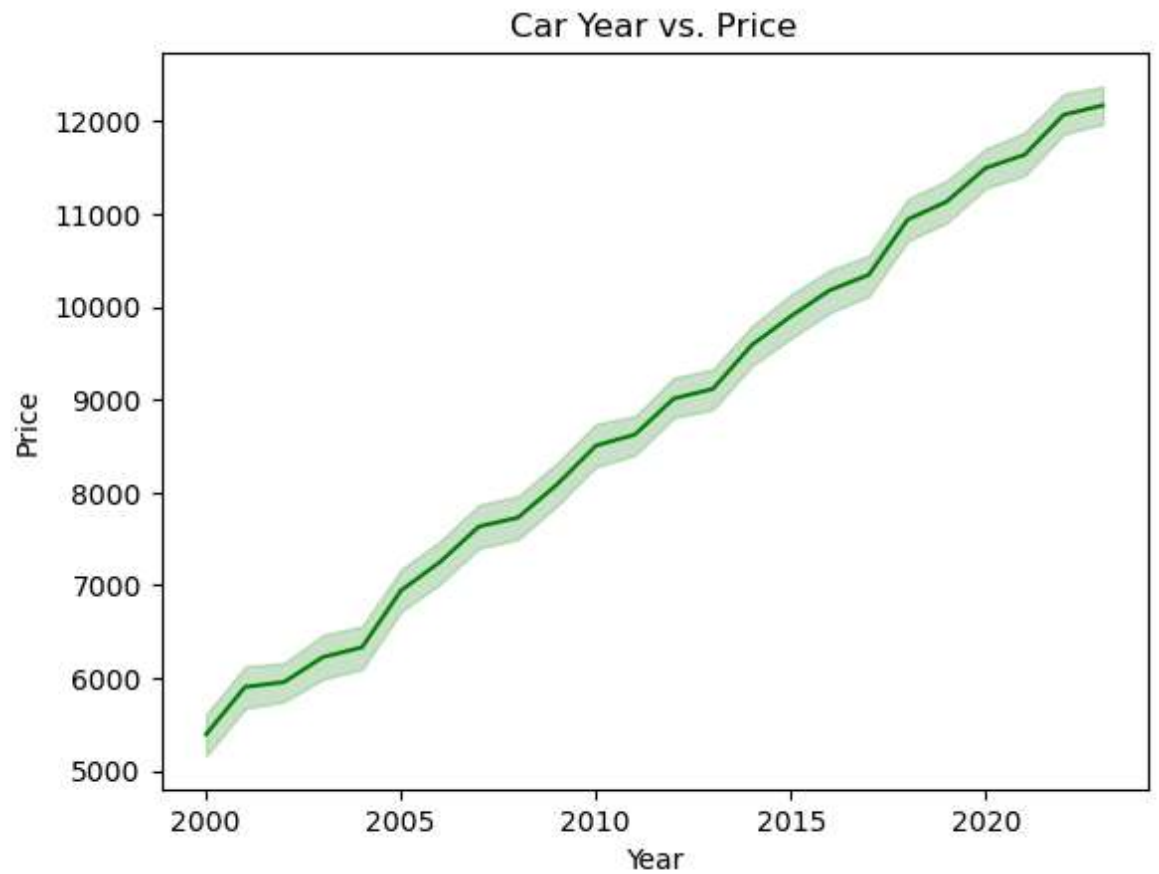
C:\Users\Saiko\AppData\Local\Temp\ipykernel\_17756\28750575.py:1: UserWarning: Ignoring `palette` because no `hue` variable has been assigned.  
sns.scatterplot(x='Engine\_Size', y='Price', data=df, palette='viridis')

Out[24]: Text(0.5, 1.0, 'Engine Size vs. Price')




```
In [25]: sns.lineplot(x='Year', y='Price', data=df, color='green')  
plt.title('Car Year vs. Price')
```

```
Out[25]: Text(0.5, 1.0, 'Car Year vs. Price')
```



```
In [26]: ▶ sns.scatterplot(x='Mileage', y='Price', data=df, color='purple', alpha=0.6)  
plt.title('Mileage vs. Price')
```

Out[26]: Text(0.5, 1.0, 'Mileage vs. Price')



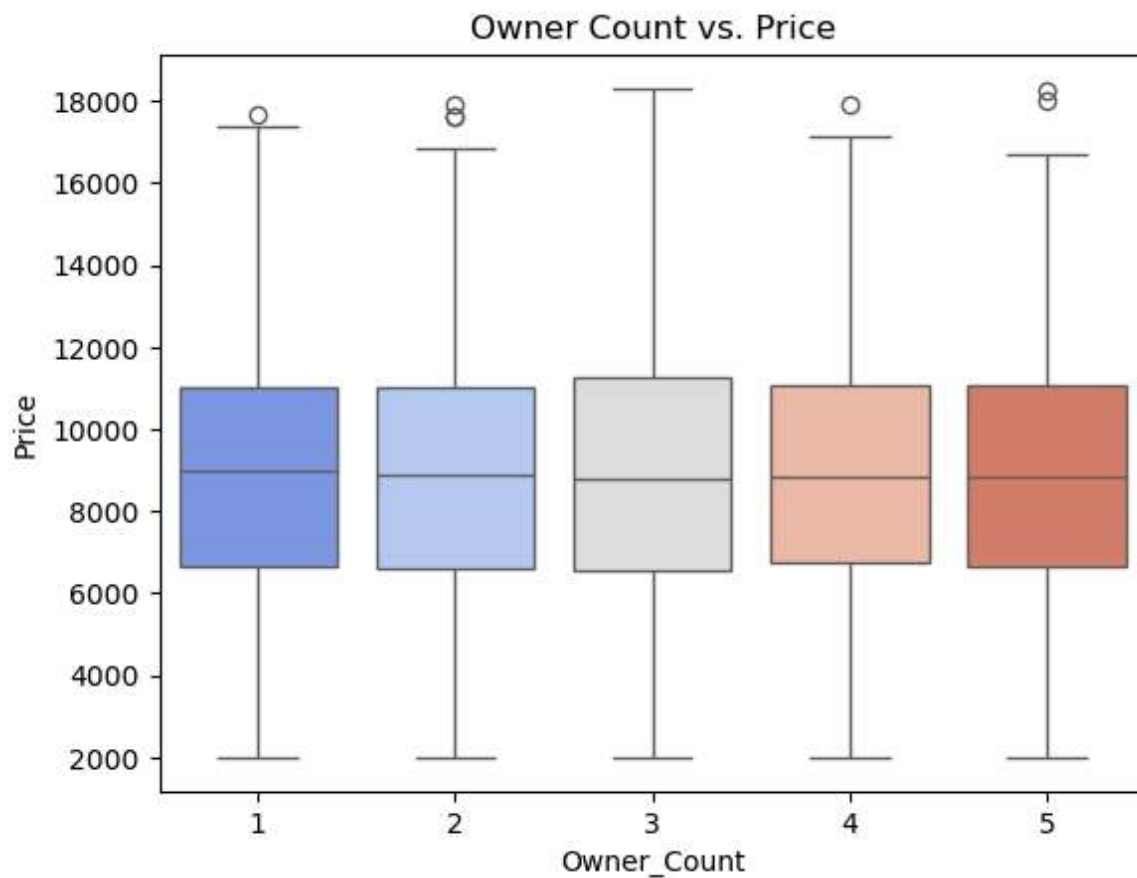
```
In [27]: sns.boxplot(x='Owner_Count', y='Price', data=df, palette='coolwarm')  
plt.title('Owner Count vs. Price')
```

C:\Users\Saiko\AppData\Local\Temp\ipykernel\_17756\1912413548.py:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Owner_Count', y='Price', data=df, palette='coolwarm')
```

Out[27]: Text(0.5, 1.0, 'Owner Count vs. Price')



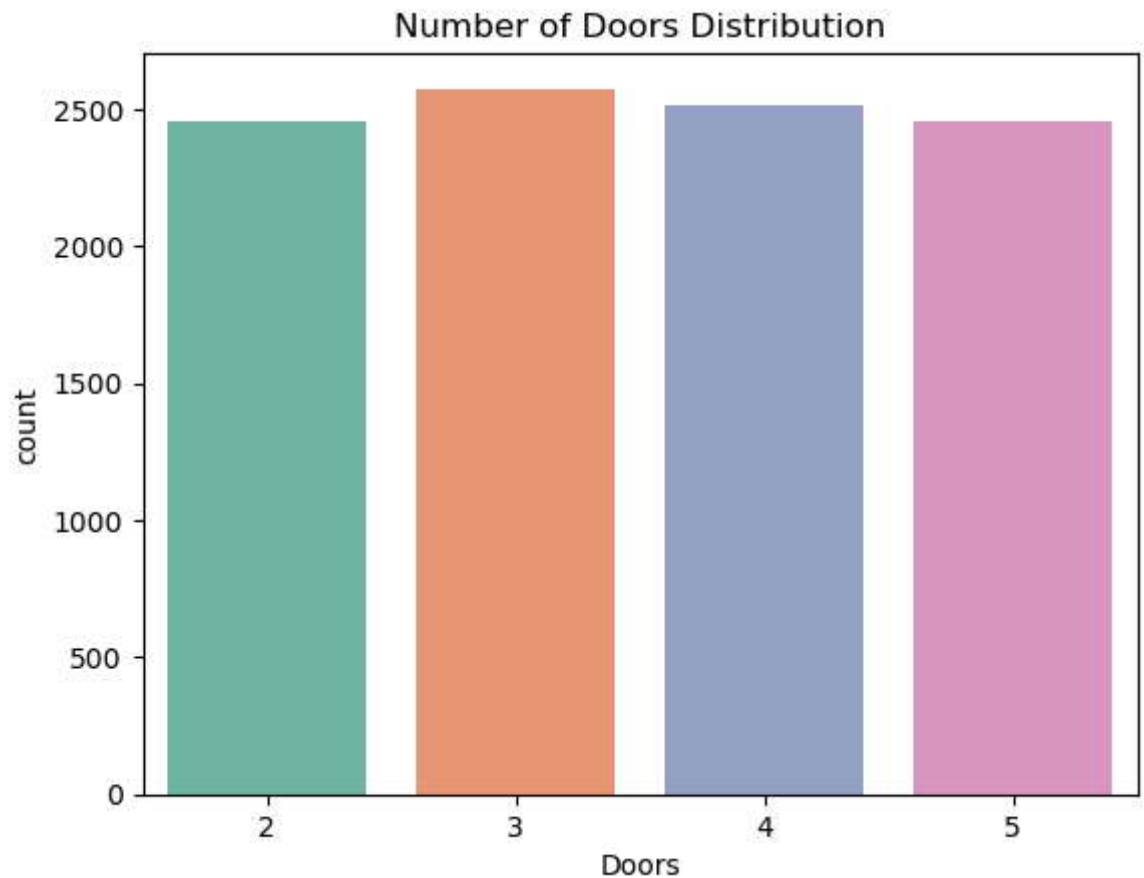
```
In [28]: sns.countplot(x='Doors', data=df, palette='Set2')  
plt.title('Number of Doors Distribution')
```

C:\Users\Saiko\AppData\Local\Temp\ipykernel\_17756\283035502.py:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Doors', data=df, palette='Set2')
```

Out[28]: Text(0.5, 1.0, 'Number of Doors Distribution')





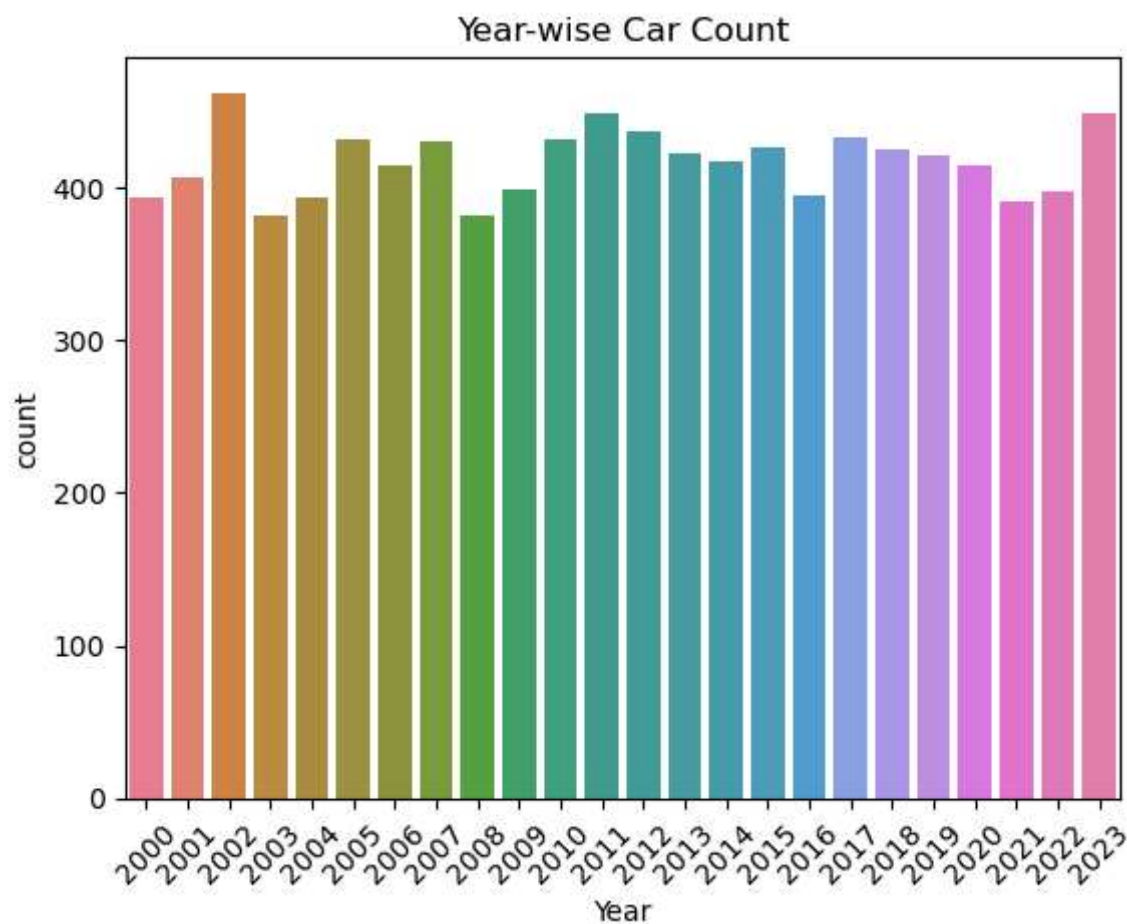
```
In [29]: sns.countplot(x='Year', data=df, palette='husl')
plt.xticks(rotation=45)
plt.title('Year-wise Car Count')
```

C:\Users\Saiko\AppData\Local\Temp\ipykernel\_17756\4274427044.py:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Year', data=df, palette='husl')
```

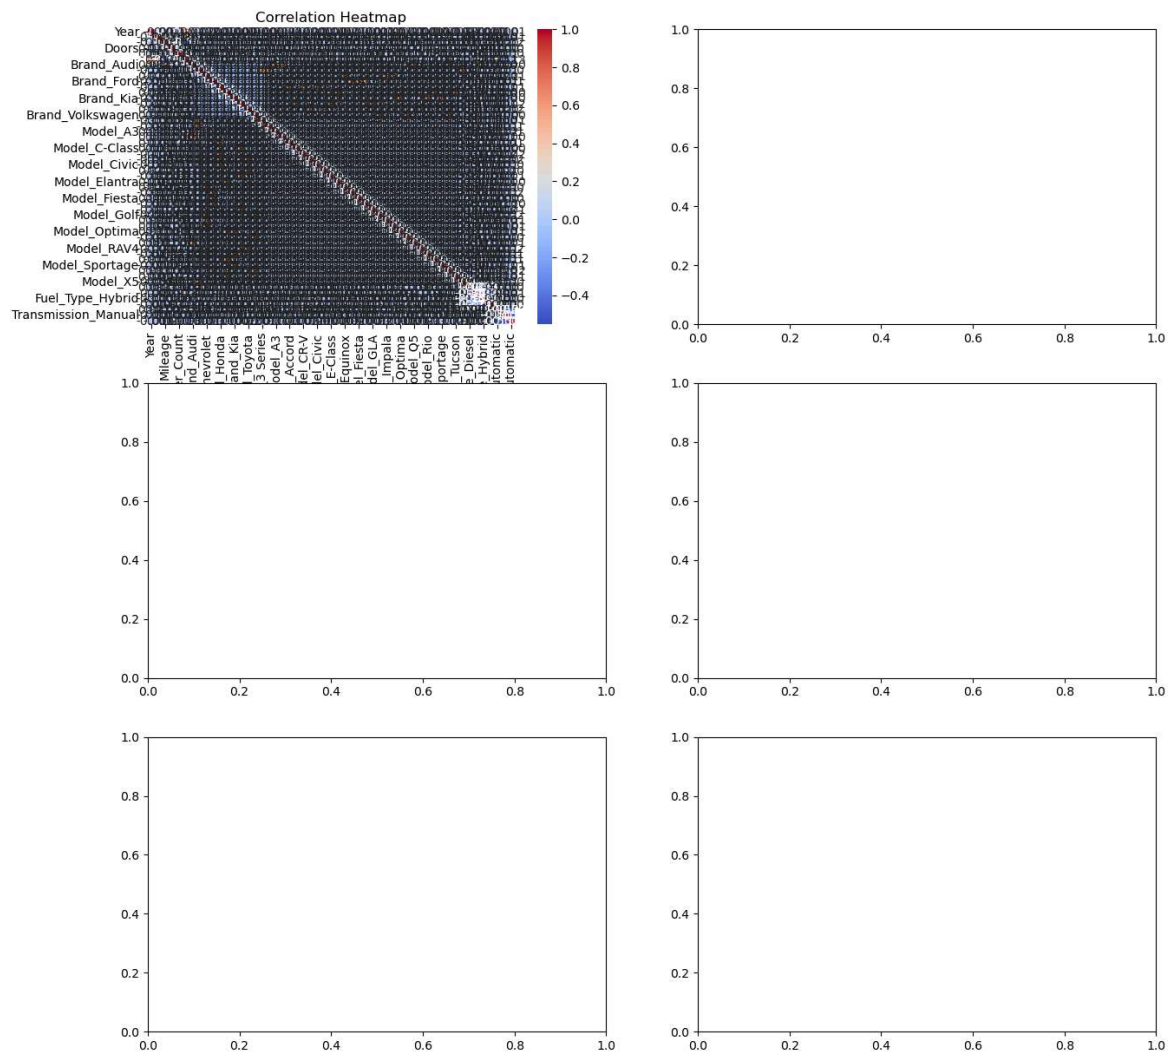
Out[29]: Text(0.5, 1.0, 'Year-wise Car Count')






```
In [30]: fig, axes = plt.subplots(3, 2, figsize=(15, 15))
axes = axes.flatten()
corr = df.corr(numeric_only=True)
sns.heatmap(corr, annot=True, fmt=".2f", cmap='coolwarm', ax=axes[0])
axes[0].set_title('Correlation Heatmap')
```

Out[30]: Text(0.5, 1.0, 'Correlation Heatmap')

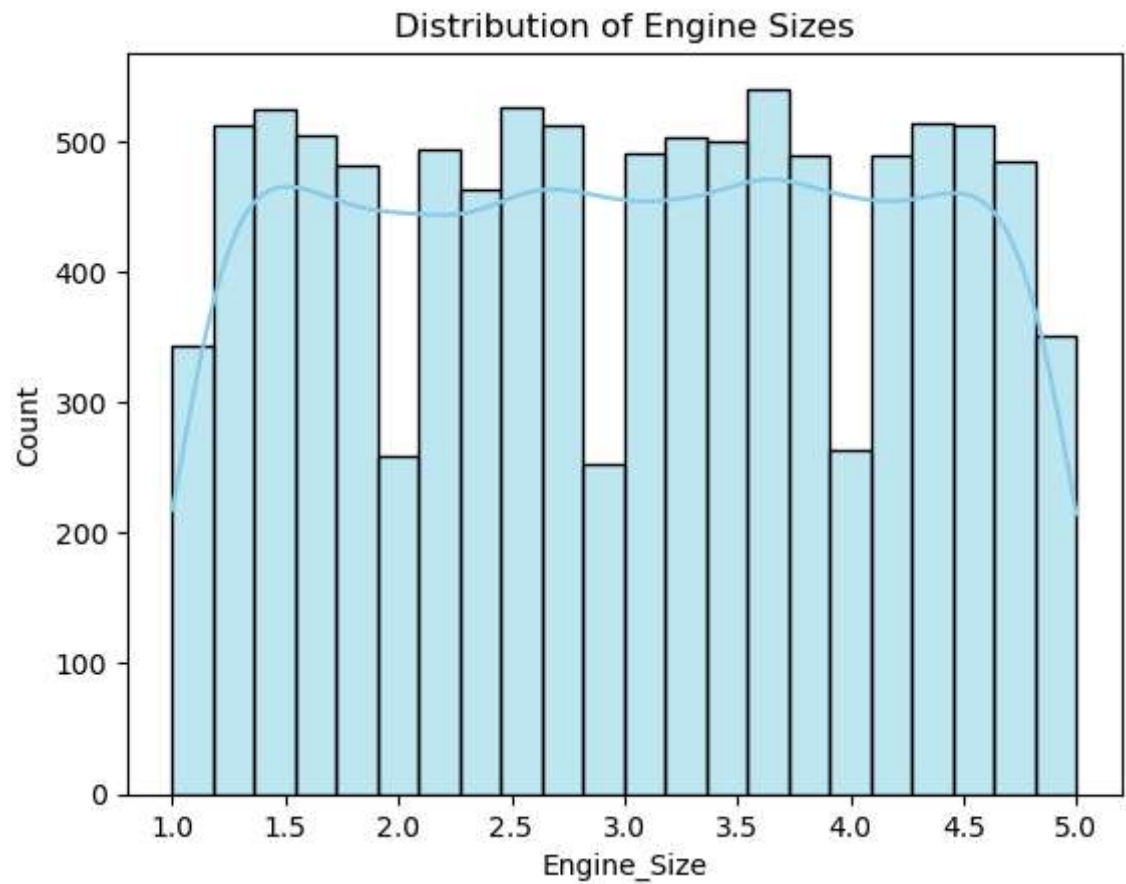


```
In [31]: ▶ fig, axes = plt.subplots(3, 2, figsize=(15, 15))  
axes = axes.flatten()
```



```
In [32]: ▶ sns.histplot(df['Engine_Size'], kde=True, color='skyblue')  
plt.title('Distribution of Engine Sizes')
```

Out[32]: Text(0.5, 1.0, 'Distribution of Engine Sizes')

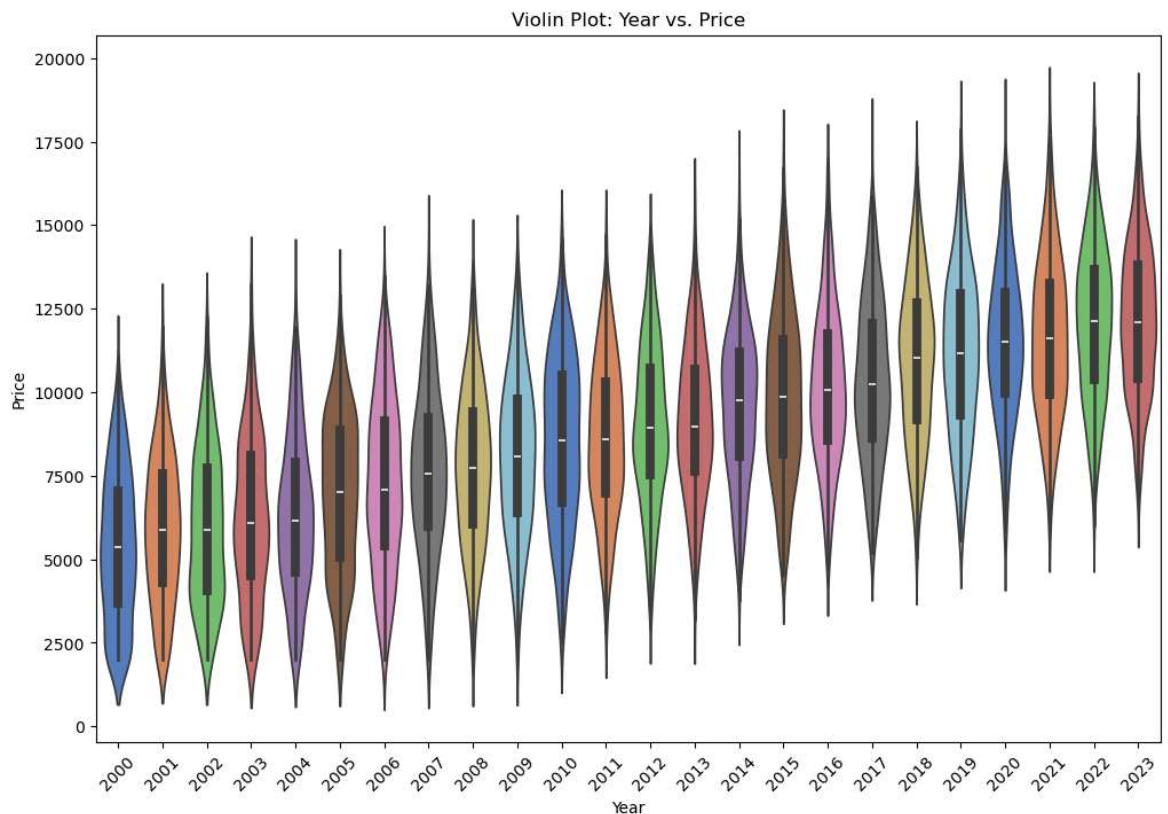


```
In [33]: plt.figure(figsize=(12, 8))
sns.violinplot(x='Year', y='Price', data=df, palette='muted')
plt.title('Violin Plot: Year vs. Price')
plt.xticks(rotation=45)
plt.show()
```

C:\Users\Saiko\AppData\Local\Temp\ipykernel\_17756\3991876073.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.violinplot(x='Year', y='Price', data=df, palette='muted')
```



### DataSet Observation: Insights from the Car Price Dataset Visualizations

#### 1. Price Trends by Year

- The Price vs. Year scatter plot and violin plot reveal that newer cars (post-2015) have significantly higher prices, while older models tend to be cheaper.
- The median price steadily increases for recent cars, reflecting their higher market value.

#### 2. Mileage and Price Relationship

- The Mileage vs. Price scatter plot highlights a clear trend:
- Cars with higher mileage tend to have lower prices, reflecting wear and depreciation.
- Lower-mileage cars retain their value better, especially newer models.

#### 3. Engine Size Impact

- The Engine Size vs. Price and Engine Size vs. Mileage scatter plots show that:

- Cars with larger engines generally have higher prices, indicating more powerful and premium vehicles.
  - Larger engines tend to have lower mileage, suggesting they may be used less frequently or preserved for special purposes.
4. Doors and Pricing The Price Distribution by Number of Doors plot shows that:
- 4-door cars dominate the market and have the widest price range, likely due to their popularity and versatility.
  - 2-door cars are generally priced lower, possibly because they are less practical for families.
  - 5-door cars exhibit moderate pricing, often associated with hatchbacks or compact SUVs.
5. Ownership Patterns The Owner Count vs. Price plot reveals that:
- Cars with fewer previous owners generally have higher prices, as they are perceived to be better maintained.
  - Cars with 3 or more owners tend to have lower prices, likely due to increased wear and potential maintenance issues.
6. Fuel Type and Transmission Insights From the earlier plots (which included fuel and transmission types):
- Hybrid and electric cars have higher prices, reflecting their growing demand and eco-friendliness.
  - Automatic cars generally have higher resale values compared to manual cars, indicating a consumer preference for convenience.
7. Yearly Car Trends The Year vs. Number of Doors and Year-wise Car Count plots show that:
- The production of cars peaked between 2015 and 2020, indicating higher availability of newer cars in the dataset.

In [ ]: ▶